

Groovy Security Executive Summary



AI-Native Governance vs Traditional DLP — Whiteout AI Capabilities Overview

Version 3.1 — June 2026

The Problem: AI Adoption Has Outpaced Security

Enterprises are rapidly adopting generative AI tools — ChatGPT, Claude, Gemini, Copilot, Cursor — across every department. Traditional Data Loss Prevention (DLP) solutions were built for a pre-AI world: file transfers, email attachments, and network perimeter enforcement. They fundamentally cannot secure **conversational data exfiltration** — the primary risk vector introduced by AI tools.

The core challenge: Across every department, employees are pasting sensitive data into public AI tools without realizing the risk:

- **HR** — candidate evaluations and compensation data
- **Finance** — forecasts and unreleased earnings
- **Legal** — privileged communications and contract terms
- **Developers** — API keys, credentials, and proprietary source code
- **Healthcare** — patient records embedded in clinical queries

Every one of these is a potential compliance violation — spanning **HIPAA, GDPR, PCI-DSS, SOX, FERPA, CCPA**, and attorney-client privilege.

Whiteout AI: Purpose-Built for the AI Era

Whiteout AI by Groovy Security is an enterprise AI governance platform that enables organizations to safely adopt generative AI while enforcing compliance, data security, and full auditability. Unlike broad-scope platforms or traditional DLP, Whiteout AI is **deployed, production-ready, and proven** — not a waitlist or roadmap promise. **No proxy, no network interference** — pure endpoint protection.

60+

BUILT-IN POLICIES

23+

AI PLATFORMS COVERED

99.19%

BENCHMARK ACCURACY

Coverage: Every AI Touchpoint, One Platform

Whiteout AI intercepts AI interactions at **every point of use** across the enterprise:

Enforcement Point	Coverage
Browser Extension	Chrome, Firefox, Edge, Safari — ChatGPT, Claude, Gemini, Copilot, Perplexity, and 23+ AI platforms
Desktop Guard	Native AI desktop apps — file drop scanning, attach interception
IDE Extensions	VS Code, Cursor, Windsurf, JetBrains — Claude Code, Copilot, Codex monitoring
MCP Governance	Controls which external tools AI agents can access (GitHub, Google Drive, Jira, etc.)
Secure Internal AI	Compliance-checked AI for sensitive queries — data never leaves the org
Infrastructure Agents	EKS, EC2, ECS, Lambda — monitor or enforce AI API calls without client deployment

Groovy Security | Public — For informational purposes only. Does not constitute a binding agreement.

Competitive Differentiation

Whiteout AI vs Traditional DLP & Emerging AI Security Platforms



Head-to-Head Comparison

Capability	Whiteout AI (Groovy Security)	Traditional DLP / Emerging Platforms
Detection Method	LLM-powered semantic analysis — understands conversational context and natural language	Regex, keyword matching, pattern rules — misses data embedded in conversation or over-blocks legitimate usage

Capability	Whiteout AI (Groovy Security)	Traditional DLP / Emerging Platforms
Accuracy	99.19% on a 15,915-prompt public benchmark	High false positive rates — tuning-intensive
Time to Deploy	Immediate — 60+ policies, zero configuration	Weeks-to-months of rule writing and tuning
Product Maturity	Deployed & Production-Ready	Many AI-specific entrants are Pre-Launch / Waitlist
Interception Points	Browser, Desktop App, IDE, MCP, Internal AI, Infrastructure	Typically network perimeter or single-channel only
AI Platform Coverage	23+ AI platforms across web, desktop, and IDE	Often requires proxy configuration, causing network overhead — no coverage for certificate-pinned desktop applications
Fail-Safe Behavior	Allows with warning on error — never blocks productivity	Blocks on error — interrupts workflows
Prompt Redaction	✓ Automatic — sanitized resubmission with smart placeholders	✗ Binary block/allow only
Secure Internal AI	✓ Built-in compliance-checked AI interface	✗ Not offered
MCP / AI Tool Governance	✓ Per-tool audit + access control for AI agents	✗ No agent-level governance
Prompt Injection Defense	✓ IDE hook-based + output scanning	~ Limited or roadmap-only
Enterprise Integration	SSO across 9 identity providers, SIEM webhooks, MDM (Intune, Jamf), 35+ integrations	Varies — often limited at early stage

Why LLM-Powered Detection Matters

Traditional Pattern Matching Misses

- "Help me draft a response about the Smith family's trust fund restructuring" (*financial PII in context*)
- "Review this discharge summary for accuracy" (*PHI embedded in clinical language*)
- "Here's our Q3 pipeline — can you help forecast?" (*unreleased financials*)

Whiteout AI Catches

- Understands that a "discharge summary" with names and diagnoses contains PHI
- Recognizes unreleased financial data regardless of how it's phrased
- Detects privileged legal content even in casual conversation

Policy Coverage by Compliance Domain

Domain	Policies	Frameworks
Protected Health Information	5 policies	HIPAA, HITECH
Personal Identifiable Information	5 policies	CCPA, PCI-DSS
GDPR & Privacy	7 policies	GDPR, LGPD, UK GDPR
Code & Intellectual Property	4 policies	Trade Secret, IP
Financial & Business	7 policies	SOX, SEC, GLB
Legal & Privilege	8 policies	Attorney-Client, Work Product
Security Operations	7 policies	NIST, ISO 27001
Confidential Business	5 policies	Trade Secrets, NDA
Education	12 policies	FERPA

Groovy Security | Public — For informational purposes only. Does not constitute a binding agreement.

Platform Architecture & Key Differentiators



GROOVY SECURITY

Architecture: Intercept at the Point of Use

Unlike network-perimeter DLP that monitors traffic after the fact, Whiteout AI intercepts **before the prompt is sent** — at the keyboard, at the send button, at the file drop. This is the only way to prevent conversational data leakage in real time.

Six Differentiators That Extend Groovy Security's Capabilities

1. Conversational DLP — Not Pattern Matching

Whiteout AI uses LLM-powered semantic analysis to understand *what data means in context*, not just what it looks like. This is the fundamental gap in every traditional DLP and most AI security startups. Our compliance engine delivers **deep contextual understanding and accuracy** across all 9 compliance domains.

2. Universal AI Coverage — Every Touchpoint

A single platform governs browser-based AI (23+ platforms), native desktop AI apps, IDE coding tools (VS Code, Cursor, JetBrains), MCP-connected AI agents, and infrastructure-level API calls. No other solution covers this breadth from a single management plane.

3. Fail-Safe Design — Security Without Productivity Loss

Compliant prompts proceed with zero perceptible latency. When network errors occur, prompts are allowed with a warning — never blocked. Blocked prompts receive **automatic redaction** with smart placeholders so users can resubmit sanitized versions immediately. Security that blocks productivity doesn't get adopted.

4. Deployed & Proven — Not a Waitlist

Whiteout AI is deployed and in production today — proven in a live clinical pilot. Many emerging AI security platforms are still in pre-launch waitlist mode with unspecified delivery timelines. When evaluating vendors for an AI Due Diligence framework, **deployed capability beats roadmap promises**.

5. Enterprise-Grade Integrations

Full SSO support across 9 identity providers (Okta, Microsoft Entra, Google Workspace, Auth0, Ping, and more) with SCIM 2.0 provisioning, SIEM webhook delivery for SOC teams, MDM integration (Intune, Jamf), and 35+ enterprise integrations (GitHub, Jira, Slack, Confluence, Notion, and more). Whiteout AI fits into existing enterprise security stacks — it doesn't require ripping and replacing.

6. AI Agent Governance (MCP)

As enterprises adopt AI agents that take actions autonomously (via Model Context Protocol), Whiteout AI provides per-tool access control and audit logging for agent interactions. This is a **forward-looking capability** that no traditional DLP addresses and few AI security platforms have considered.

Ideal Customer Profiles

Regulated Industries

- **Healthcare** — HIPAA, PHI protection
- **Financial Services** — SOX, PCI-DSS, SEC compliance
- **Legal** — Attorney-client privilege protection
- **Education** — FERPA, student data (12 policies)
- **Government** — Security clearance, NIST alignment

Enterprise Use Cases

- **Shadow AI Governance** — visibility into unauthorized AI usage
- **IP Protection** — prevent source code and trade secrets in prompts
- **Developer Workflows** — IDE governance for AI coding tools
- **AI Due Diligence** — audit trail and policy enforcement for AI adoption
- **Compliance Reporting** — exportable audit logs and analytics

Groovy Security | Public — For informational purposes only. Does not constitute a binding agreement.

Positioning Within an AI Due Diligence Framework



How Whiteout AI Enhances Partner-Led AI Governance Programs

The AI Due Diligence Umbrella

Organizations pursuing AI adoption need a layered due diligence approach. Whiteout AI provides the **real-time enforcement and audit layer** that complements advisory, assessment, and posture management capabilities. For partners building an AI Due Diligence practice,

Whiteout AI delivers the technical enforcement that turns policy recommendations into measurable, auditable controls.

Due Diligence Layer	Capability	Whiteout AI Role
Policy Definition	Define what data is sensitive and what usage is acceptable	60+ pre-built policies + custom policy creation per group
Real-Time Enforcement	Prevent violations before they happen	Browser, desktop, IDE, MCP interception with LLM analysis
Audit & Compliance	Prove governance is working to auditors and regulators	Immutable audit trail, SIEM integration, exportable reports
Visibility & Analytics	Understand how AI is being used across the organization	Admin dashboard with usage stats, risk trends, coverage metrics
Incident Response	Detect and respond to policy violations	Real-time SOC webhooks, risk scoring, violation alerts
Access Governance	Control which AI tools and integrations are permitted	AI app allow/block lists, MCP tool permissions, LLM access controls

Why This Matters for AI Due Diligence Practices

Groovy Security doesn't compete with advisory and assessment firms — we complement them. When a partner identifies that a client needs AI governance, Whiteout AI provides the deployed technology layer that makes governance enforceable. The result: a complete AI Due Diligence offering with advisory expertise backed by real-time technical controls.

What Partners Bring

- Client relationships and trust
- AI risk assessment expertise
- Regulatory compliance advisory
- Due diligence framework definition
- Ongoing governance consulting

What Whiteout AI Adds

- Real-time policy enforcement technology
- 60+ policies across 9 compliance domains
- Immutable audit trail for compliance proof
- Universal AI platform coverage
- Enterprise integration (SSO, SIEM, MDM)

Groovy Security Product Suite

Whiteout AI

Enterprise AI governance platform providing real-time compliance enforcement, audit, and analytics across every AI touchpoint — browser, desktop, IDE, MCP, and infrastructure. 60+ built-in policies across 9 compliance domains with LLM-powered semantic analysis at 99.19% benchmark accuracy. Includes admin dashboard, secure internal AI, SSO/SIEM/MDM integrations, and full audit trail.

Maestro

AI-driven automated penetration testing platform that deploys **18 specialized AI agents** with **203 MCP tools** to find vulnerabilities across web apps, APIs, cloud, Kubernetes, and identity providers — then validates them through real red team-style exploitation, proving actual impact, not just scanner output. Built on Kali Linux and deployed locally so all data stays on your network.

- **209-Test Assessment Matrix** — 73 DAST tests, 60 identity tests, 29 cloud security tests, 24 SAST tests, 15 cross-validation tests, and 8 chain analysis tests for consistent, deterministic coverage
- **Red Team Exploitation** — goes beyond detection by testing and validating every finding through controlled exploitation to confirm real-world exploitability
- **Identity & Cloud Red Teaming** — Active Directory, Entra ID, M365, Okta, Google Workspace, and Ping attack paths, plus AWS/Azure/GCP and Kubernetes exploitation
- **Multi-Track DAST + SAST** — parallel dynamic, static, cloud, and identity analysis with cross-validation and source code context enrichment
- **Chain Analysis** — links individual findings into compound exploit paths with severity recalculation and defense-in-depth analysis
- **Evidence-Based Reporting** — HTML, PDF, and Markdown with executive summaries, technical breakdowns, reproduction steps, and prioritized remediation
- **Compliance Mapping** — automated mapping to OWASP Top 10, OWASP API Top 10, NIST 800-53, PCI-DSS, and CWE standards with CVSS v3.1 scoring

Ready to explore a partnership?

partners@groovysec.com | groovysec.com